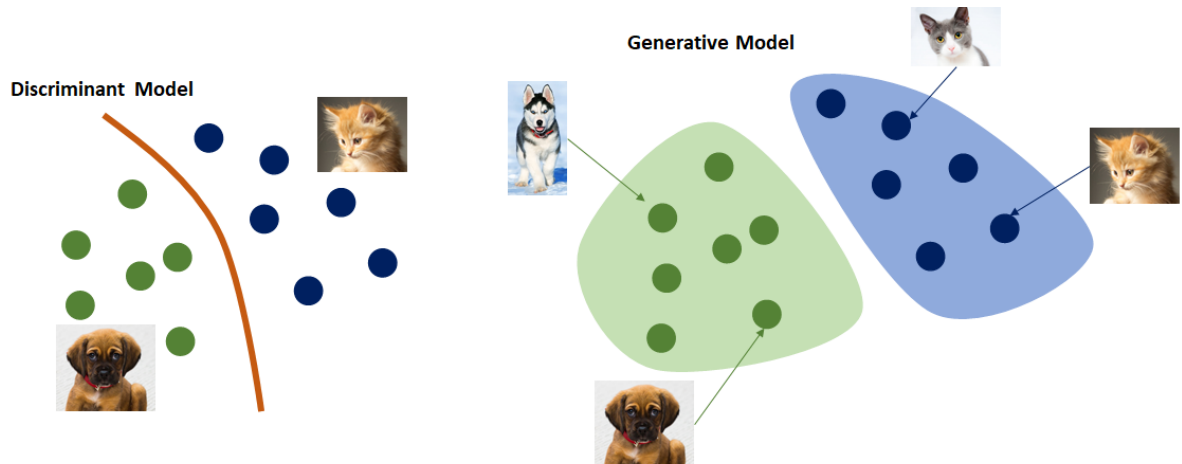




EE 046746 - Technion - Computer Vision

Tutorial 02 - Probabilistic Discriminative Learning

Elias Nehme



- [Image source](#)



Agenda

- [Machine Learning Overview](#)
 - [What is Machine Learning?](#)
 - [Supervised Learning](#)
- [Discriminative vs Generative Models](#)
 - [Discriminative Modelling](#)
 - [Generative Modelling](#)
- [Probabilistic Discriminative Models](#)
 - [General Idea](#)
 - [Binary Classification Example](#)
- [Recommended Videos](#)
- [Credits](#)

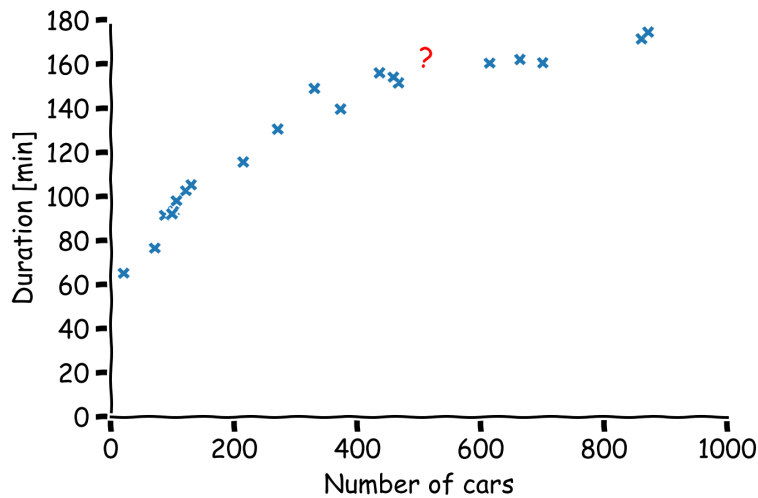


Machine Learning Overview




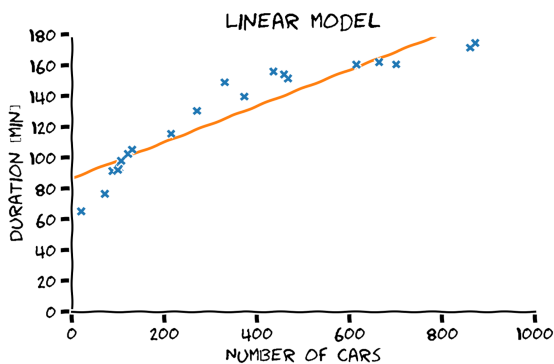


What is Machine Learning?

- Wikipedia definition: "the study of computer algorithms that can improve automatically through experience and by the use of data"
- What does that really mean? Let's see an example.
- **Example 1:** Consider the task of predicting the duration of travelling the road given the current amount of cars.



- Naturally one can choose multiple models to predict the duration.
- For example, we can limit ourselves to parametric models (eg polynomials)
- Which of the following models should we choose?



- Higher polynomial degree ensures better fit to data.

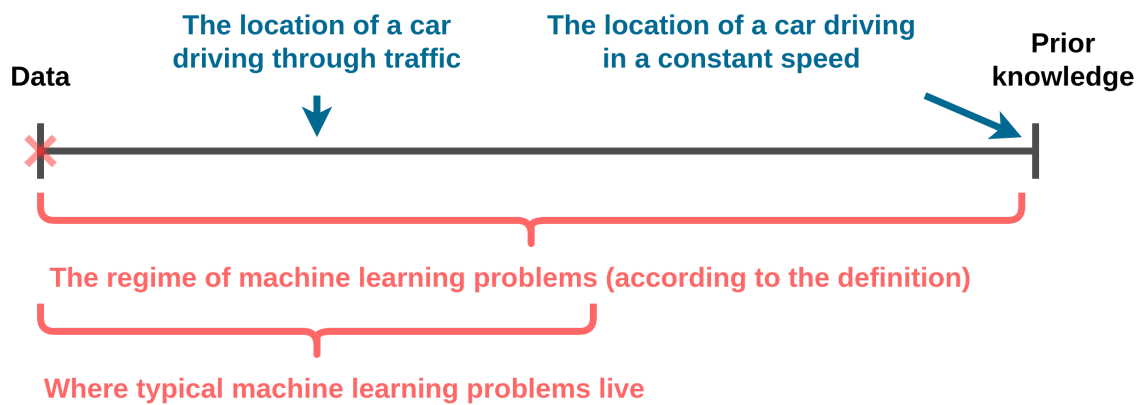
- However, better fit to data does not necessarily ensure better "generalization" error.
- Assuming we have prior knowledge the function is smooth we might choose a 4th order polynomial.




Data vs Prior Knowledge

- If we have a solid understanding of the system (geometry, physics, etc.), we rely mainly on prior knowledge.
- However, if we do not have such understanding, we try to infer the system based on collected data.
- For example, coming back to the example from earlier, if the car is driving at constant speed: duration = distance/speed.


What should the model be based on?



- **Example 2:** Consider the task of classifying a credit card deal being legit or fraud based on some characteristics like the Price, and the physical distance of the store from the card holder's home.



- We would like to create a prediction function that tells us based on these 2 features whether a credit card deal is legit or fraud.
- For example here a naive implementation of such a function.





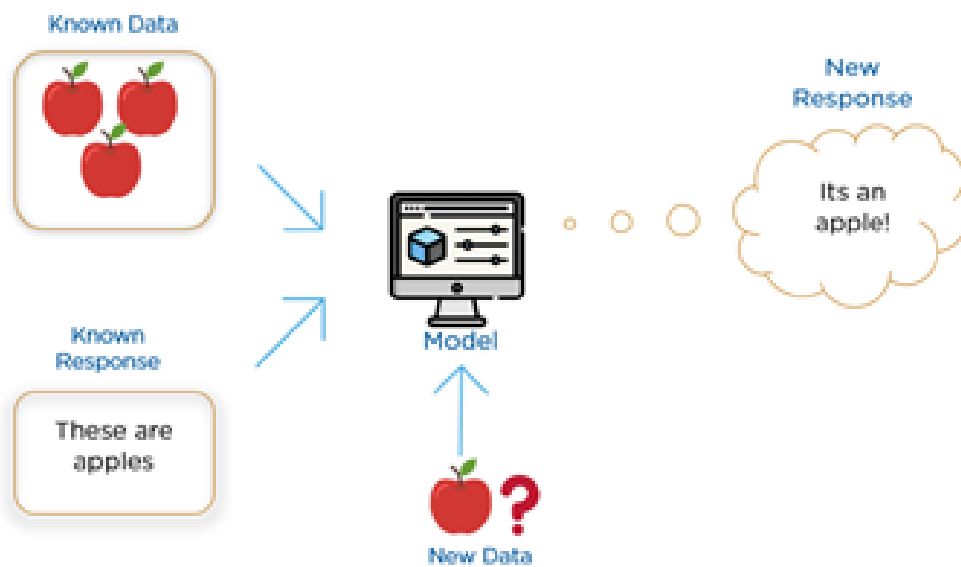
Types of ML Setups

- There are 3 main different types of machine learning setups:
 - **Supervised Learning** (This tutorial)
 - Unsupervised Learning
 - Reinforcement Learning




What is Supervised Learning?

- Simplest form of machine learning (easiest to understand). Data is given in the form of examples with labels.
- Algorithm is "trained" to predict the label for each example, while being given feedback for its answers.
- when fully-trained the learning algorithm will be able to observe a new, never-before-seen example and predict a good label for it.
- For example, label emails as spam, classify images as apples, classify cats vs dogs, predict location of tumor in medical images etc.



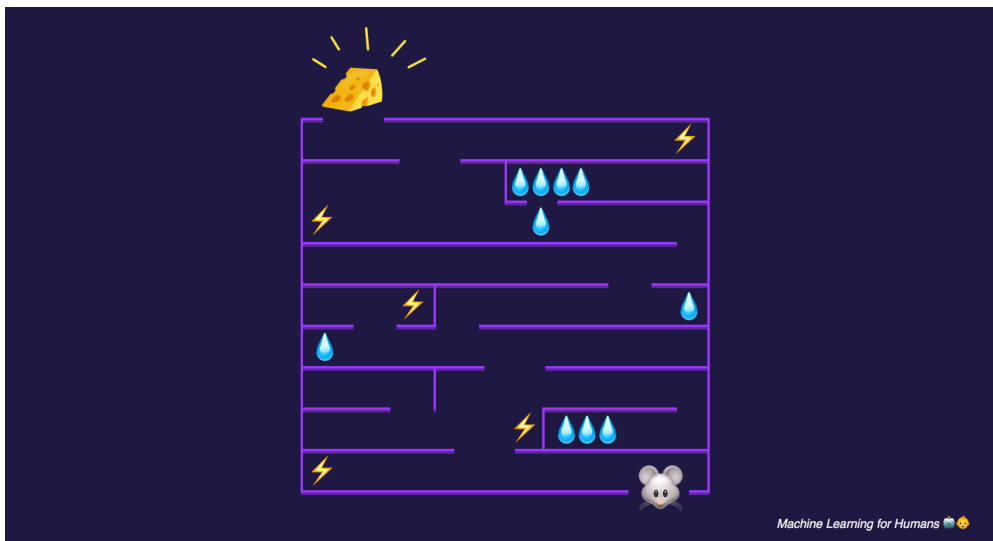
What is Unsupervised Learning?

- In this case we only have examples (data) with no labels.
- Algorithm is "trained" to understand the properties of the data.
- After training, it can learn to group, cluster, and/or organize the data in a way such that a human (or other intelligent algorithm) can come in and make sense of the newly organized data. (Example in lecture on K-means).



What is Reinforcement Learning?

- Fairly different from the previous two approaches, and beyond the scope of this course.
- Essentially it tries to learn from mistakes with sparse/not-frequent supervision.
- For example, learning a "policy" to navigate a maze with obstacles.



Supervised Learning

- The basis of all other ML problems. Relates to estimation problems from statistical theory.
- The estimation problem is the following: We want to predict the value of unknown random variable (y) based off other known random variables (x).
- Usually in statistics, we assume the distribution of all random variables is known.
- In Supervised learning, we assume we only have a finite sample from this distribution.
- Therefore, our estimator will be built based off the finite sample only.



Supervised Learning - Notation

- Labels y - The random variable we are trying to predict
- Observations/Measurements x - the random variables we are basing our predictor on.
- Predictor/Estimator $\hat{y} = h(x)$ - is the prediction function.
- Dataset $\mathcal{D} = \{x^{(i)}, y^{(i)}\}_{i=1}^N$ - comprised of N pairs of i.i.d samples from the joint distribution.




Supervised Learning - Problem Types

- Depending on the values that y can take we classify into 2 sub-types:
- Continuous labels y - a regression problem (estimating travel duration example).
- Discrete labels y - a classification problem (predicting legit vs fraud transaction example).




Supervised Learning - Regression

- Regression problem: x = Number of cars, y = Duration time:



Supervised Learning - Classification

- Classification problem: $x = [\text{Price}, \text{Distance from home}]^T$, $y \in \{\text{Legit}, \text{Fraud}\}$:



Supervised Learning - Optimal Estimator

- Any function mapping $h : x \rightarrow y$ is a valid estimator.
- Optimally we would like our estimator to make no mistakes.
- Due to the labels (y) being a random variable, this is impossible.
- Therefore, we need a way to compare the errors of different estimators to choose the best one.



Supervised Learning - General Solution Paradigm

- Define a mathematical criterion that measures how good an estimator is doing
- Choose a large enough family of models such that one of them will be good enough
- Search all models in the chosen family for the best one.
- (Easier said than done..)



Supervised Learning - Cost Function

- $C(h)$ Gives each estimator a score, lower score = better estimator.
- Optimal estimator $h^*(x)$ is the one with the minimal score:

$$h^* = \underset{h}{\operatorname{argmin}} C(h)$$

- Usually chosen from a subset of widely used functions.



Supervised Learning - Loss and Risk Functions

- The loss function ℓ gives a score for a single prediction:

$$\ell(h(x), y) = \ell(\hat{y}, y)$$

- The cost is then defined as the expectation over the joint distribution:

$$C(h) = \mathbb{E} [\ell(h(x), y)]$$

- A familiar synonym for the cost function is the "Risk" function:

$$R(h) \equiv C(h)$$



Supervised Learning - Popular Loss/Risk Functions

- Classification Problems - Misclassification Rate:

$$\ell(\hat{y}, y) = I \{ \hat{y} \neq y \}, \quad R(h) = \mathbb{E} [I \{ h(x) \neq y \}]$$

- Optimal Classifier given by the Mode:

$$h^*(x) = \underset{y}{\operatorname{argmax}} p(y|\mathbf{x} = x)$$

- Regression Problems - Mean Squared Error:

$$\ell(\hat{y}, y) = (\hat{y} - y)^2, \quad R(h) = \mathbb{E} [(h(x) - y)^2]$$

- Optimal Regressor given by the Conditional Mean:

$$h^*(x) = \mathbb{E} [y|\mathbf{x} = x]$$



Supervised Learning - Unknown Distribution and Empirical Risk

- Problem: We don't actually have access to the posterior distribution $p(y|x)$
- Solution 1 (Generative Models): Estimate the joint distribution based on the dataset

$$\mathcal{D} = \{x^{(i)}, y^{(i)}\}_{i=1}^N$$

- Solution 2 (Discriminative Models): Estimate the expectation empirically based on the dataset $\mathcal{D} = \{x^{(i)}, y^{(i)}\}_{i=1}^N$
- That is, replace the expectation with the empirical mean:

$$R(h) = \mathbb{E} [\ell(h(x), y)] \approx \hat{R}(h) = \hat{\mathbb{E}}_{\mathcal{D}} [\ell(h(x), y)] = \frac{1}{N} \sum_{i=1}^N \ell(h(x^{(i)}), y^{(i)})$$

- Expected to converge in probability when $N \rightarrow \infty$
- Using the empirical risk introduces the problem of overfitting (e.g. polynomial fitting example).



Supervised Learning - Performance Evaluation

- Goal: Learn a model from the given dataset that performs well on **unseen** data.
- For that we would like to evaluate our estimator's performance on data not used in training.
- This is usually achieved by splitting the dataset into two:
 - Training Set - $\mathcal{D}_{\text{train}}$ - used to build our estimator $h^*(x)$
 - Test Set - $\mathcal{D}_{\text{test}}$ - used to evaluate performance on new data
- Performance on the test set can be approximated by the empirical mean:

$$\text{test score} = \frac{1}{N_{\text{test}}} \sum_{x^{(i)}, y^{(i)} \in \mathcal{D}_{\text{test}}} \ell(h^*(x^{(i)}), y^{(i)})$$



Supervised Learning - Parametric Models

- Usually, it is more practical to learn a parametric estimator $h(x; \theta)$ with parameters θ , than searching the entire space of functions for a general $h(x)$.
- This has the benefit of simplifying the optimization and usually reduces overfitting.
- Examples of Parametric Functions:
 - Linear functions: $h(x; \theta) = \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3$
 - Polynomials: $h(x; \theta) = \theta_1 + \theta_2 x_1 + \theta_3 x_1^2 + \theta_4 x_1^3$
 - **Neural networks!** (Next Week)



Supervised Learning - Parametric Models Cont.

- Note that finding the best $h(x; \theta)$ is now broadcasted to finding the optimal set of parameters θ .
- This translates the minimization of the risk functions from earlier to rely on θ :

$$h^* = \underset{h}{\operatorname{argmin}} R(h) \rightarrow \theta^* = \underset{\theta}{\operatorname{argmin}} R(h(x; \theta))$$

- For the empirical risk case, this is approximated by:

$$\theta^* = \underset{\theta}{\operatorname{argmin}} \frac{1}{N} \sum_{i=1}^N \ell(h(x^{(i)}; \theta), y^{(i)})$$




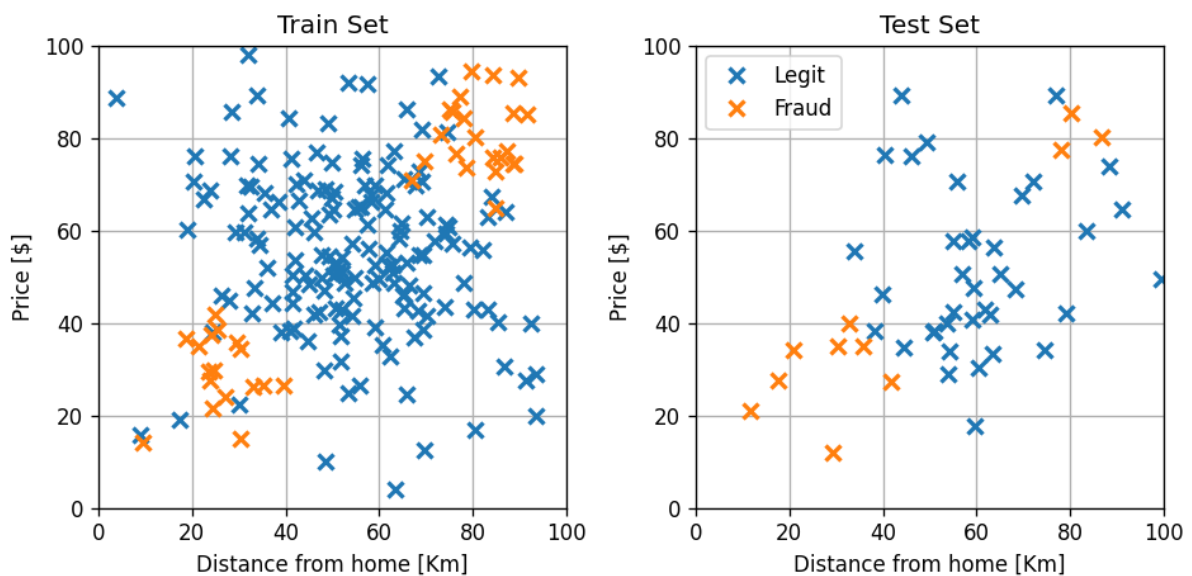
Discriminative vs Generative Models

- Recall: the problem is we don't have the posterior probability distribution $p(y|x)$
- Solution 1 (Generative Models): Estimate the joint distribution based on the dataset
 $\mathcal{D} = \{x^{(i)}, y^{(i)}\}_{i=1}^N$
- Solution 2 (Discriminative Models): Estimate the expectation empirically based on the dataset
 $\mathcal{D} = \{x^{(i)}, y^{(i)}\}_{i=1}^N$



Classification Example

- Let us get back to the classification example, and denote Legit as $y = 0$, and Fraud as $y = 1$.
- First, we split the data into 80% training and 20% test:




Discriminative Modelling

- Main Idea - Build an estimator/classifier/discriminator from the empirical risk directly:

$$\hat{h}_{\mathcal{D}}(x) = \underset{h}{\operatorname{argmin}} \hat{\mathbb{E}}_{\mathcal{D}} [\ell(h(x), y)] \rightarrow \hat{h}_{\mathcal{D}}(x) = \underset{h}{\operatorname{argmin}} \frac{1}{N} \sum_{i=1}^N \ell(h(x^{(i)}), y^{(i)})$$

- The estimator can be non-parametric ($h(x)$) or parameteric ($h(x; \theta)$).
- Here, we will discuss a non-parametric estimator (Nearest Neighbor)
- In the lecture you will cover also a parametric estimator (Support Vector Machines).




Discriminative Modelling - Nearest Neighbor Classifier

- Algorithm That classifies based on the label of the nearest sample in the dataset:
- We are given a dataset $\mathcal{D} = \{x^{(i)}, y^{(i)}\}_{i=1}^N$, and test sample x_q that we want to predict the label for.
- The algorithm is extremely simple and comprised of two steps:
 - Find the the index of the nearest sample in the dataset:

$$i = \underset{i}{\operatorname{argmin}} \|x_q - x^{(i)}\|_2$$



- Predict the label to be the label at the found index:

$$\hat{y} = y^{(i)}$$



Discriminative Modelling - Nearest Neighbor Classifier Cont.



- Here are the resulting decision boundaries on the training set for the classification problem:



Discriminative Modelling - Nearest Neighbor Classifier Cont.

- The resulting error on the test score is 12%:



$$\text{test score} = \frac{1}{N_{\text{test}}} \sum_{x^{(i)}, y^{(i)} \in \mathcal{D}_{\text{test}}} \ell(h^*(x^{(i)}), y^{(i)}) = 0.12$$



Discriminative Modelling - Nearest Neighbor Classifier Cont.


- The resulting test score can be reduced to 10% using a modified algorithm with $K = 5$ neighbors (Details in the lecture):

$$\text{test score} = \frac{1}{N_{\text{test}}} \sum_{x^{(i)}, y^{(i)} \in \mathcal{D}_{\text{test}}} \ell(h^*(x^{(i)}), y^{(i)}) = 0.1$$



Discriminative Modelling - Popular Algorithms


- K Nearest Neighbors
- Decision Trees
- Support Vector Machines
- Linear Least Squares Regression



Generative Modelling


- First: Estimate the unknown joint distribution $\hat{p}(x, y)$: Conveniently separated to $\hat{p}(x|y)$ and $\hat{p}(y)$
- Second: Employ Bayes rule to get the approximated posterior (up to normalization):

$$\hat{p}(y|x) = \frac{\hat{p}(x|y)\hat{p}(y)}{p(x)} = \frac{\hat{p}(x|y)\hat{p}(y)}{\int_y \hat{p}(x|y)\hat{p}(y)}$$
- Third: Use the approximated posterior to get the final classifier: $h^*(x) = \underset{y}{\operatorname{argmax}} \hat{p}(y|x = x)$
- The estimation of $p(x|y)$ can be non-parametric (KDE) or parameteric (eg Gaussians).
- Here, we will discuss a parametric estimator for $p(x|y; \theta)$ (Quadratic Discriminant Analysis)
- In the lecture you will cover a non-parametric estimator $p(x|y)$ (Naive Bayes Classifier).




Generative Modelling - Quadratic Discriminant Analysis


- Assume we have a mixed joint probability distribution $p(x, y)$, where some of the variables are continuous x , and some are discrete y .
- It is convenient in this case to write down $p(x, y) = p(x|y)p(y)$ and estimate each part separately:
 - $p(y)$ can be simply estimated using label proportions in the dataset
 - $p(x|y)$ can be estimated separately for each value of y
- If we further assume a parametric form $p(x|y; \theta)$, we can turn the problem into estimating θ per y



Generative Modelling - Quadratic Discriminant Analysis Cont.



- Back to the example of the credit card transactions:
- Training set have 200 data points with 160 Legit ($y = 0$) and 40 Fraud ($y = 1$).
- Therefore $p(y)$ can be estimated using the proportions: $p(y = 0) = \frac{160}{200} = 0.8$ and $p(y = 1) = \frac{40}{200} = 0.2$





Generative Modelling - Quadratic Discriminant Analysis Cont.

- Next, we estimate a conditional density $p(x|y)$ for each y separately: $p(x|y = 0)$ and $p(x|y = 1)$.
- This can be done with a non-parametric method (eg KDE) or a parametric one (eg Gaussian density).
- For the parametric case, we can derive the optimal parameters using Maximum Likelihood Estimation (MLE).



Generative Modelling - Quadratic Discriminant Analysis Cont.

- Assuming a parametric Gaussian density, then for each $p(x|y)$ we will estimate an expectation vector $\mu_{x|y}$ and a covariance matrix $\Sigma_{x|y}$.
- The resulting estimated posterior probabilities are given by:

$$\hat{p}(y = 0|x) = \frac{\hat{p}(x|y = 0)\hat{p}(y = 0)}{p(x)}, \quad \hat{p}(y = 1|x) = \frac{\hat{p}(x|y = 1)\hat{p}(y = 1)}{p(x)}$$


- Finally the resulting estimator is given by the mode of the resulting posterior using Bayes rule:

$$h^*(x) = \underset{y}{\operatorname{argmax}} \hat{p}(y|\mathbf{x} = x) = \underset{y}{\operatorname{argmax}} \{\hat{p}(y = 0|x), \hat{p}(y = 1|x)\}$$

- This translates to the criterion:

$$h^*(x) = \begin{cases} 0 & \text{if } \hat{p}(x|y=0)\hat{p}(y=0) > \hat{p}(x|y=1)\hat{p}(y=1) \\ 1 & \text{otherwise.} \end{cases}$$



- Substituting the Gaussian distribution in $p(x|y = 0/1)$ results in quadratic boundaries, hence the name.



Generative Modelling - Quadratic Discriminant Analysis Cont.

- Note that the orange Gaussian is probably a poor model choice and a mixture model of two Gaussians would have worked much better.
- The resulting test score is 8%:


$$\text{test score} = \frac{1}{N_{\text{test}}} \sum_{x^{(i)}, y^{(i)} \in \mathcal{D}_{\text{test}}} \ell(h^*(x^{(i)}), y^{(i)}) = 0.08$$



Generative Modelling - Popular Algorithms

- Linear Discriminant Analysis
- Quadratic Discriminant analysis
- Naive Bayes

- Markov Random Fields



Generative Modelling - Limitations

- Suffers from the curse of dimensionality:
 - Space coverage becomes increasingly more difficult for higher dimensions of x
 - Number of samples n needed to estimate the conditional density $p(x|y)$ is exponential in the dimension: $\approx n^d$
 - Naive solution via assuming independence (Naive Bayes Classifier)
- Models we can work with are very limited, because we need to satisfy:
 - $p(x, y; \theta) \geq 0, \forall x, y, \theta$
 - $\int \int p(x, y; \theta) = 1, \forall \theta$



Probabilistic Discriminative Models

- Recall: the problem is we don't have the posterior probability distribution $p(y|x)$
- Solution 1 (Generative Models): Estimate the **joint** distribution based on the dataset $\mathcal{D} = \{x^{(i)}, y^{(i)}\}_{i=1}^N$
- Solution 2 (Discriminative Models): Estimate the expectation empirically based on the dataset $\mathcal{D} = \{x^{(i)}, y^{(i)}\}_{i=1}^N$
- **Solution 3 (Prob. Discriminative Models):** Estimate the **posterior** distribution **directly** based on the dataset $\mathcal{D} = \{x^{(i)}, y^{(i)}\}_{i=1}^N$



Probabilistic Discriminative Models -General Idea

- Usually these methods in most books will be called simply discriminative without the "probabilistic".
- However, most good books will still differentiate between different types of discriminative models although without special names.
- Here, we will estimate directly $p(y|x)$, usually by parameterizing it to $p(y|x; \theta)$.
- Parameter estimation as usual will be done with maximum likelihood.



Probabilistic Discriminative Models -General Idea Cont.

- For classification problems, the function $p(y|x; \theta)$ should satisfy:
 - $p(y|x; \theta) \geq 0, \forall x, y, \theta$
 - $\sum_{y=1}^C p(y|x; \theta) = 1, \forall x, \theta$

- The second here is much simpler than the condition for generative models which was:
 $\int \int p(x, y; \theta) = 1, \forall \theta$
- Such models can be easily constructed, for example for $C = 2$ classes (binary classification), we only demand:

$$p(y = 0|x; \theta) + p(y = 1|x; \theta) = 1 \forall x, \theta$$

- Any parametric function $f(x; \theta)$ that returns values between 0 and 1 can define a valid model like so:


$$\begin{aligned} p(y = 1|x) &= f(x; \theta) \\ p(y = 0|x) &= 1 - f(x; \theta) \end{aligned}$$



Probabilistic Discriminative Models - Binary Classification Example

- Assume we are dealing with a binary classification problem $y \in \{0, 1\}$.
- Using the sigmoid function we can even relax the range condition of $f(x; \theta)$.
- Known as Logistic Regression in the literature.
- Any parametric model of the following form is valid:

$$\begin{aligned} p(y = 1|x) &= \sigma(f(x; \theta)) \\ p(y = 0|x) &= 1 - \sigma(f(x; \theta)) \end{aligned}$$



Probabilistic Discriminative Models - Binary Classification Example

- Estimating the parameters θ with MLE we get the following objective function:

$$\begin{aligned}
\theta^* &= \underset{\theta}{\operatorname{argmin}} - \sum_{i=1}^N \log(p(y^{(i)}|x^{(i)}; \theta)) \\
&= \underset{\theta}{\operatorname{argmin}} - \sum_{i=1}^N I\{y^{(i)} = 1\} \log(\sigma(f(x^{(i)}; \theta))) + I\{y^{(i)} = 0\} \log(1 - \sigma(f(x^{(i)}; \theta))) \\
&= \underset{\theta}{\operatorname{argmin}} - \sum_{i=1}^N y^{(i)} \log(\sigma(f(x^{(i)}; \theta))) + (1 - y^{(i)}) \log(1 - \sigma(f(x^{(i)}; \theta)))
\end{aligned}$$


- Resulting estimator after optimization:

$$h(x) = \underset{y}{\operatorname{argmax}} p(y|x; \theta) = \begin{cases} 1 & \text{if } \sigma(f(x; \theta)) > 0.5 \\ 0 & \text{otherwise.} \end{cases} = \begin{cases} 1 & \text{if } f(x; \theta) > 0 \\ 0 & \text{otherwise.} \end{cases}$$



Probabilistic Discriminative Models - Binary Classification Example


- Coming back to the fraud example from earlier (taking only one part of the space for simplicity):



Probabilistic Discriminative Models - Binary Classification Example

- Fitting a linear parametric function: $f(x; \theta) = \theta^T x$
- The resulting test score is 2%:

$$\text{test score} = \frac{1}{N_{\text{test}}} \sum_{x^{(i)}, y^{(i)} \in \mathcal{D}_{\text{test}}} \ell(h^*(x^{(i)}), y^{(i)}) = 0.02$$



Probabilistic Discriminative Models - Binary Classification Example

- For $f(x; \theta)$ taken as a second order polynomial, the resulting test score is 0% (no mistakes!):

$$\text{test score} = \frac{1}{N_{\text{test}}} \sum_{x^{(i)}, y^{(i)} \in \mathcal{D}_{\text{test}}} \ell(h^*(x^{(i)}), y^{(i)}) = 0$$



Recommended Videos



Warning!

- These videos do not replace the lectures and tutorials.
- Please use these to get a better understanding of the material, and not as an alternative to the written material.

Video By Subject

- Machine Learning Course by Andrew Ng - [Coursera](#)
- Lectures 1-7, EE 046195 by Omer Yair - [Technion](#)



Credits

- EE 046746 Spring 2021 - [Elias Nehme](#)
- Lectures 1-7, EE 046195 Spring 2021 - [Omer Yair](#)
- What are the Types of Machine Learning? - [Hunter Heidenreich](#)
- Icons from [Icon8.com](#) - <https://icons8.com>