



EE 046746 - Technion - Computer Vision

Elias Nehme

Tutorial 07 - Homography, Alignment & Panoramas



Agenda

- Matching Local Features
- Parametric Transformations
- Computing Parametric Transformations
 - Affine
 - Projective
- RANSAC
- Panorama
 - Warping
 - Image Blending (Feathering)
- Kornia & Transformations in Deep Learning
- Recommended Videos
- Credits

The largest panorama in the world (2014): Mont Blanc

In2WHITE Video

In2WHITE Full Image

Homographic usage examples:

<http://blog.flickr.net/en/2010/01/27/a-look-into-the-past/>

<https://www.instagram.com/albumplusart/>

In [33]:

```
%%html
<iframe src="http://www.in2white.com/" width="700" height="600"></iframe>
```



Filippo B. (<http://www.filippobengini.com>)

▲ Sponsored by: ▲

©2015

In [29]:

```
# imports for the tutorial
import numpy as np,sys
import matplotlib.pyplot as plt
from PIL import Image
import cv2
from scipy import signal
%matplotlib inline
```

In [30]:

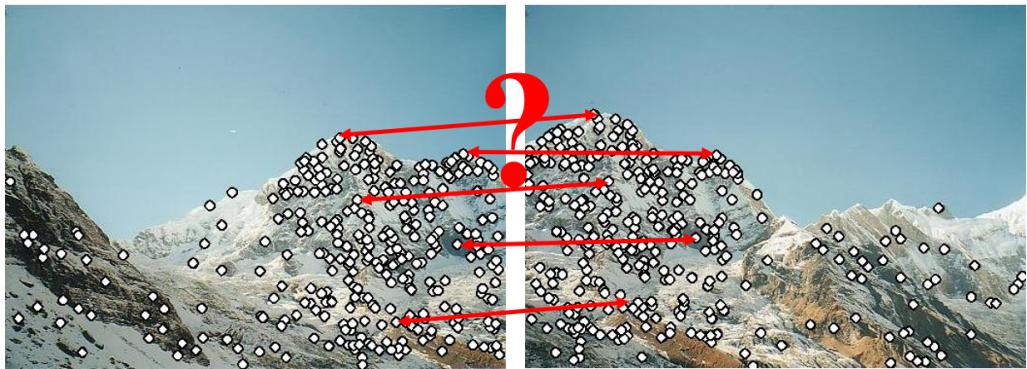
```
# plot images function
def plot_images(image_list, title_list, subplot_shape=(1,1), axis='off', fontsize=30, figsize=(4,4), c
    plt.figure(figsize=figsize)
    for ii, im in enumerate(image_list):
        c_title = title_list[ii]
        if len(cmap) > 1:
            c_cmap = cmap[ii]
        else:
            c_cmap = cmap[0]
        plt.subplot(subplot_shape[0], subplot_shape[1],ii+1)
        plt.imshow(im, cmap=c_cmap)
        plt.title(c_title, fontsize=fontsize)
        plt.axis(axis)
```



Matching Local Features

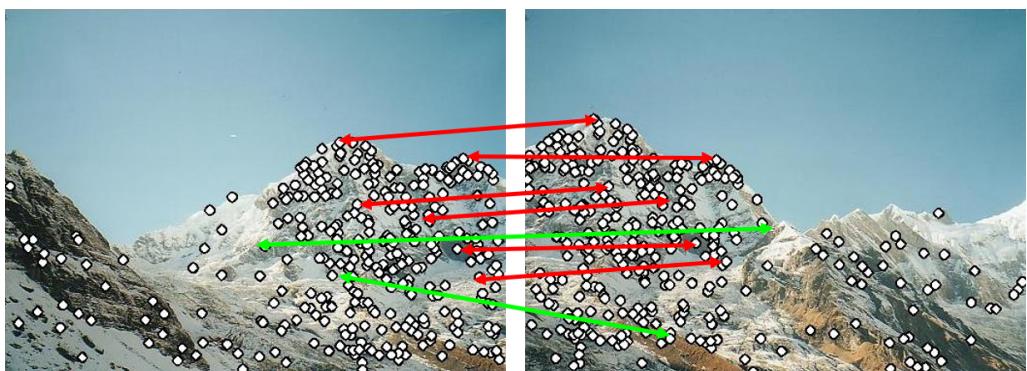
Feature matching

- We know how to detect and describe good points
- Next question: How to match them?



Typical feature matching results

- Some matches are correct
- Some matches are incorrect



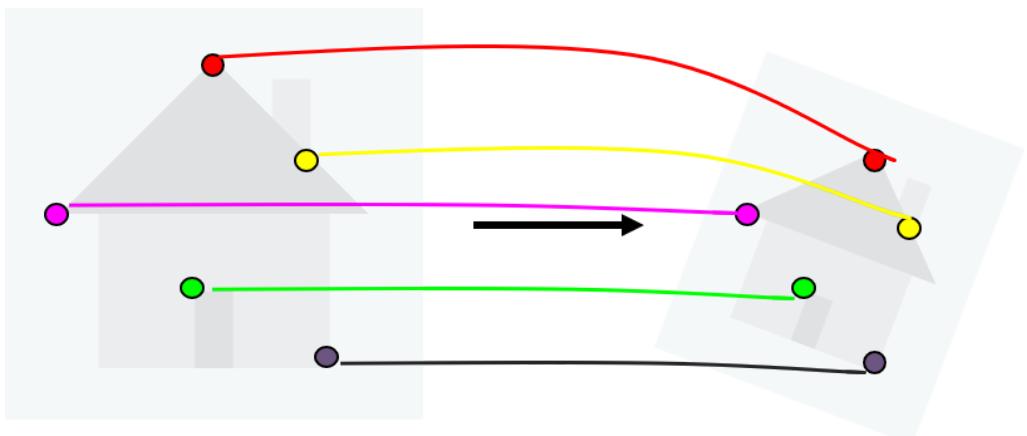
- Solution: search for a set of geometrically consistent matches



Parametric Transformations

Image Alignment

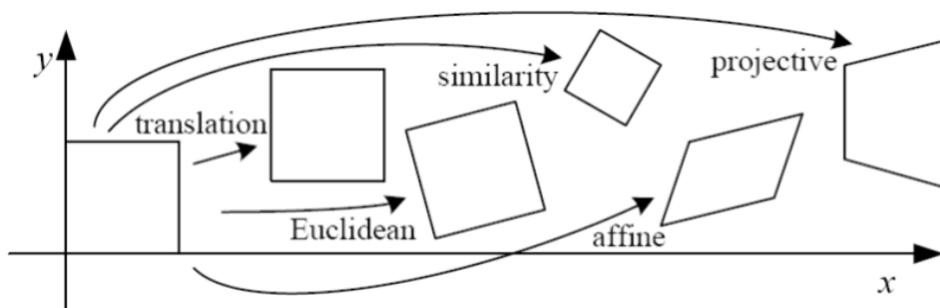
Given a set of matches, what parametric model describes a geometrically consistent transformation?



Basic 2D Transformations

Transformation	Matrix	# DoF	Preserves	Icon
translation	$\begin{bmatrix} I & t \end{bmatrix}_{2 \times 3}$	2	orientation	
rigid (Euclidean)	$\begin{bmatrix} R & t \end{bmatrix}_{2 \times 3}$	3	lengths	
similarity	$\begin{bmatrix} sR & t \end{bmatrix}_{2 \times 3}$	4	angles	
affine	$\begin{bmatrix} A \end{bmatrix}_{2 \times 3}$	6	parallelism	
projective	$\begin{bmatrix} \tilde{H} \end{bmatrix}_{3 \times 3}$	8	straight lines	

(x is in homogeneous coordinates)



Parametric (Global) Warping

Examples of parametric warps:

```
In [31]: im = cv2.imread('./assets/tut_8_exm.jpg')
im = cv2.cvtColor(im, cv2.COLOR_BGR2RGB)
rows, cols, d = im.shape
print(im.shape)

# Translation:
tx,ty = [10,20]
h_T = np.float32([[1,0,tx],[0,1,ty],[0,0,1]])

# Rotation
theta = np.deg2rad(20)
h_R = np.float32([[np.cos(theta), -np.sin(theta),0],[np.sin(theta),np.cos(theta),0],[0,0,1]])
# Affine
h_AFF = np.array([[ 1.26666667,-0.5,-60],[-0.33333333,1.66666667],[0,0,1]])

H = [h_T,h_R,h_AFF]

img = [im]
for h in H:
    img.append(cv2.warpPerspective(im, h,(cols,rows)))
Titles = ['Original','Translation','Rotation','Affine']

# Perspective
pts1 = np.float32([[100,77],[320,105],[100,150],[385,170]])
pts2 = np.float32([[0,0],[300,0],[0,100],[300,50]])
M = cv2.getPerspectiveTransform(pts1,pts2)
im_perspective = cv2.warpPerspective(im,M,(400,300))

# Cylindrical (non Linear)
def cylindricalWarp(img, K):
    """This function returns the cylindrical warp for a given image and intrinsics matrix K"""
    h_,w_ = img.shape[:2]
```

```

# pixel coordinates
y_i, x_i = np.indices((h_,w_))
X = np.stack([x_i,y_i,np.ones_like(x_i)],axis=-1).reshape(h_*w_,3) # to homog
Kinv = np.linalg.inv(K)
X = Kinv.dot(X.T).T # normalized coords
# calculate cylindrical coords (sin\theta, h, cos\theta)
A = np.stack([np.sin(X[:,0]),X[:,1],np.cos(X[:,0])],axis=-1).reshape(w_*h_,3)
B = K.dot(A.T).T # project back to image-pixels plane
# back from homog coords
B = B[:, :-1] / B[:, [-1]]
# make sure warp coords only within image bounds
B[(B[:,0] < 0) | (B[:,0] >= w_) | (B[:,1] < 0) | (B[:,1] >= h_)] = -1
B = B.reshape(h_,w_,-1)

img_rgba = cv2.cvtColor(img, cv2.COLOR_BGR2RGBA) #BGR2BGRA for transparent borders...
# warp the image according to cylindrical coords
return cv2.remap(img_rgba, B[:, :, 0].astype(np.float32), B[:, :, 1].astype(np.float32),
                 cv2.INTER_AREA, borderMode=cv2.BORDER_TRANSPARENT)
# Camera parameters:
K = np.array([[200,0,cols/2],[0,400,rows/2],[0,0,1]]) # mock intrinsics
img_cyl = cylindricalWarp(im, K)

```

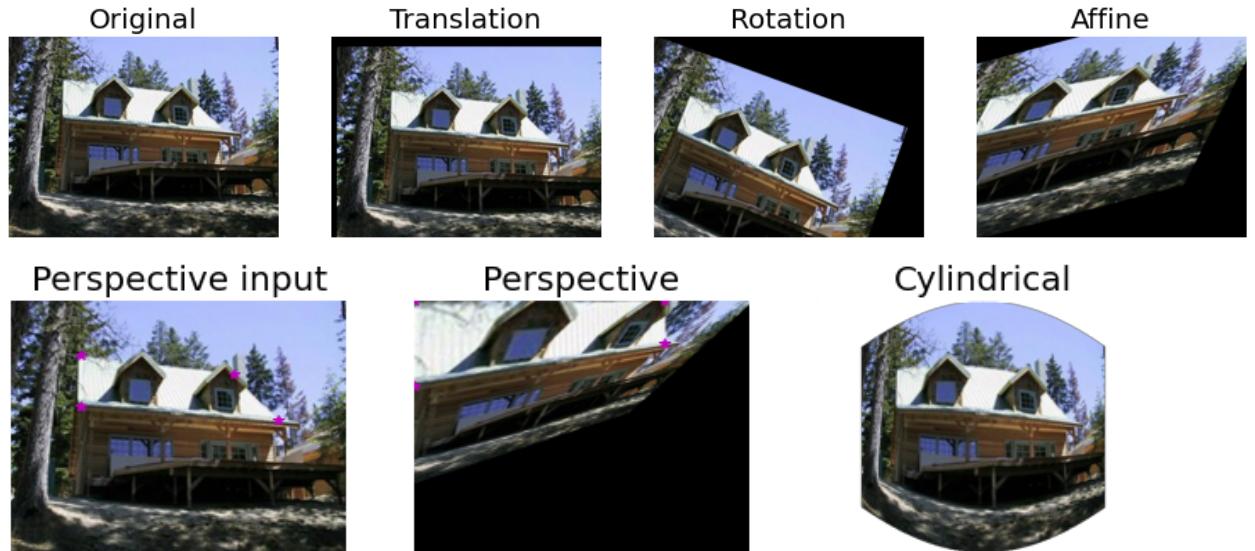
(362, 484, 3)

In [32]: plot_images(img, Titles, (1,4), figsize=(16,4), fontsize=20)

```

plt.figure(figsize=(12,4))
plt.subplot(131)
plt.imshow(im, plt.title('Perspective input', fontsize=20))
plt.plot(pts1[:,0],pts1[:,1], 'm*')
plt.axis('off')
plt.subplot(132), plt.imshow(im_perspective), plt.title('Perspective', fontsize=20)
plt.plot(pts2[:,0],pts2[:,1], 'm*')
plt.axis('off')
plt.subplot(133), plt.imshow(img_cyl)
plt.title('Cylindrical', fontsize=20)
_ = plt.axis('off')

```



- [Code source](#) - OpenCV
- [Cylinder code source](#) - More Than Technical
- [Cylinder coordinates](#)



Basic 2D Transformations

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \quad \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \quad \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos\Theta & -\sin\Theta & 0 \\ \sin\Theta & \cos\Theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translate

Scale

Rotate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Affine

$$\begin{bmatrix} x' \\ y' \\ w \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Projective
(Homography)

Basic 2D Transformations - Translation

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x + t_x \\ y + t_y \\ 1 \end{bmatrix}$$

```
In [11]: plot_images([img[0],img[1]],[Titles[0],Titles[1]],(1,2),figsize=(12,12),fontsize=20)
```

Original



Translation



Basic 2D Transformation - Rotation

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos(\theta) & -\sin(\theta) & 0 \\ \sin(\theta) & \cos(\theta) & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- Around which point do we rotate the image?

```
In [12]: plot_images([img[0],img[2]],[Titles[0],Titles[2]],(1,2),figsize=(12,12),fontsize=20)
```

Original



Rotation



Basic 2D Transformations – Translation and Rotation (2D rigid body motion)

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos(\theta) & -\sin(\theta) & t_x \\ \sin(\theta) & \cos(\theta) & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- Euclidean distances are preserved
- Combination of rotation and translation, which one applied first?

In [13]:

```
H = h_T @ h_R
im_temp = cv2.warpPerspective(im, H, (cols, rows))
plot_images([img[0], im_temp], [Titles[0], 'Rigid'], (1,2), figsize=(12,12), fontsize=20)
```

Original



Rigid



Basic 2D Transformations – Scale

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s & 0 & 0 \\ 0 & s & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

In [14]:

```
s = 1.5
H_s = np.array([[s, 0, 0], [0, s, 0], [0, 0, 1]], np.float32)
im_temp = cv2.warpPerspective(im, H_s, (cols, rows))
plot_images([img[0], im_temp], [Titles[0], 'Scale'], (1, 2), figsize=(12,12), fontsize=20)
```

Original



Scale



Basic 2D Transformations – Similarity

Similarity transform (4 DoF) = translation + rotation + scale

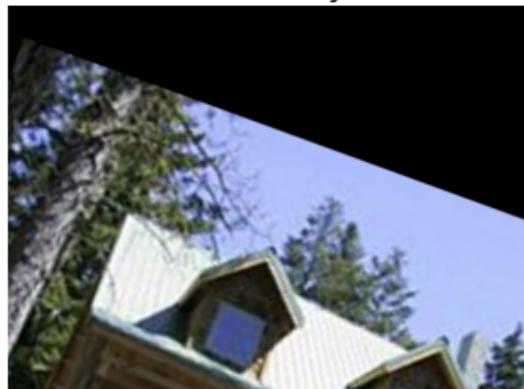
In [15]:

```
H = H_s @ h_T @ h_R
im_temp = cv2.warpPerspective(im, H, (cols, rows))
plot_images([img[0], im_temp], [Titles[0], 'Similarity'], (1,2), figsize=(12,12), fontsize=20)
```

Original



Similarity



Basic 2D Transformation - Aspect Ratio

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & 0 & 0 \\ 0 & \frac{1}{a} & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

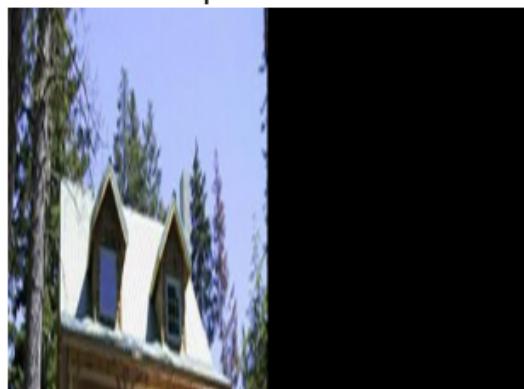
In [16]:

```
a = 1 / 2
H_a = np.array([[a, 0, 0], [0, 1 / a, 0], [0, 0, 1]], np.float32)
im_temp = cv2.warpPerspective(im, H_a, (cols,rows))
plot_images([img[0], im_temp], [Titles[0], 'Aspect Ratio'], (1,2), figsize=(12,12), fontsize=20)
```

Original



Aspect Ratio



Basic 2D Transformations – Shear

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & a & 0 \\ b & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

```
In [17]:  
a,b = (0.5, 0.1)  
H_sh = np.array([[1, a, 0],[b, 1, 0], [0, 0, 1]], np.float32)  
im_temp = cv2.warpPerspective(im, H_sh, (cols, rows))  
plot_images([img[0], im_temp], [Titles[0], 'Shear'], (1,2), figsize=(12,12), fontsize=20)
```

Original



Shear



Basic 2D Transformations – Affine

Similarity transform (6 DoF) = translation + rotation + scale + aspect ratio +shear

```
In [18]:  
H = H_a @ H_sh @ H_s @ h_T @ h_R  
im_temp = cv2.warpPerspective(im, H, (cols * 4, rows * 4))  
plot_images([img[0], im_temp], [Titles[0], 'Affine'], (1,2), figsize=(12,12), fontsize=20)
```

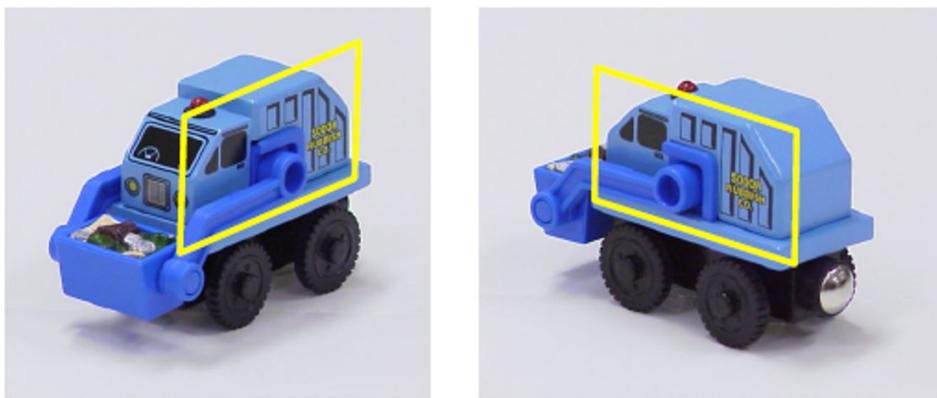
Original



Affine



- Simple fitting procedure (linear least squares)
- Approximates viewpoint change for roughly planar objects and roughly orthographic camera
- Can be used to initialize fitting for more complex models



Basic 2D Transformations – Projective - a.k.a - Homographic

$$\begin{bmatrix} u \\ v \\ w \end{bmatrix} = \begin{bmatrix} h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 \\ \textcolor{red}{h_7} & \textcolor{red}{h_8} & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

$$x' = u/w$$

$$y' = v/w$$

Non-linear!

```
In [19]: plot_images([img[0], im_perspective], [Titles[0], 'Projective'], (1, 2), figsize=(12,12), fontsize=20)
```



When do we get Homography?

Homography maps between:

- points on a plane in the world and their positions in an image
- points in two different images of the **same** plane
- two images of a 3D object where the camera has rotated but not translated

For far away objects:

- works fine for small viewpoint changes



Computing Parametric Transformations

- [Affine](#)
- [Projective](#)

Computing Affine Transformation

- Assuming we know correspondences, how do we get transformation?

$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} = \begin{bmatrix} h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix} = \begin{bmatrix} & & \dots & & & \\ x_i & y_i & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_i & y_i & 1 \\ & & & \dots & & \end{bmatrix} \begin{bmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \\ h_5 \\ h_6 \end{bmatrix}$$

$$b = Ah$$

- Solve with Least-squares $\|Ah - b\|^2$

$$h = (A^T A)^{-1} A^T b$$

In Python:

```
A_inv = pinv(A)

h = np.linalg.pinv(A)@b
```

- How many matches (correspondence pairs) do we need to solve?
- Once we have solved for the parameters, how do we compute the coordinates of the cooresponding point for any pixel (x_{new}, y_{new}) ?

$$H = \begin{bmatrix} h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 \\ 0 & 0 & 1 \end{bmatrix}$$

$$x'_{new} = Hx_{new}$$

Computing Projective Transformation

- Recall working with homogenous coordinates

$$\begin{bmatrix} u_i \\ v_i \\ w_i \end{bmatrix} = \begin{bmatrix} h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 \\ h_7 & h_8 & 1 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$

$$x'_i = u_i/w_i$$

$$y'_i = v_i/w_i$$

- We get the following non-linear equation:

$$x'_i = \frac{h_1x_i + h_2y_i + h_3}{h_7x_i + h_8y_i + h_9}$$

$$y'_i = \frac{h_4x_i + h_5y_i + h_6}{h_7x_i + h_8y_i + h_9}$$

- We can re-arrange the equation

$$\begin{bmatrix} x_i & y_i & 1 & 0 & 0 & 0 & \dots & -x_i x'_i & -y_i x'_i & -x'_i \\ 0 & 0 & 0 & x_i & y_i & 1 & \dots & -x_i y'_i & -y_i y'_i & -y'_i \end{bmatrix} \begin{bmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \\ h_5 \\ h_6 \\ h_7 \\ h_8 \\ h_9 \end{bmatrix} = \begin{bmatrix} \dots \\ 0 \\ 0 \\ \dots \end{bmatrix}$$

- We want to find a vector h satisfying

$$Ah = 0$$

where A is full rank. We are obviously not interested in the trivial solution $h = 0$ hence we add the constraint

$$\|h\| = 1$$

- Thus, we get the homogeneous Least square equation:

$$\arg \min_h \|Ah\|_2^2, \text{ s.t. } \|h\|_2^2 = 1$$

Compute Projective transformation using SVD:

$$\arg \min_h \|Ah\|_2^2, \text{ s.t. } \|h\|_2^2 = 1$$

- Let decompose A using SVD: $A = UDV^T$, where U and V are orthonormal matrix, and D is a diagonal matrix.
 - Need a reminder on SVD? [Click Here](#)
- From orthonormality of U and V :

$$\|UDV^T h\| = \|DV^T h\|$$

$$\|V^T h\| = \|h\|$$

Hence, we get the following minimization problem:

$$\arg \min_h \|DV^T h\| \text{ s.t. } \|V^T h\| = 1$$

- Substitute $y = V^T h$:

$$\arg \min_h \|Dy\| \text{ s.t. } \|y\| = 1$$

- D is a diagonal matrix with decreasing values. Then, it is clear that $y = [0, 0, \dots, 1]^T$.
- Therefore, choosing h to be the last column in V will minimize the equation.

In Python:

```
(U,D,Vh) = np.linalg.svd(A, False)
```

```
h = Vh.T[:, -1]
```

Some more options to find h :

- Lagrange multipliers - [Least-squares Solution of Homogeneous Equations](#)
- Using EVD (eigenvalue decomposition) on $A^T A$.
- If we know our transformation is nearly Affine we can get an approximate solution using linear least squares

RANSAC

- [Tutorial 2](#)
- [Lecture 7](#)

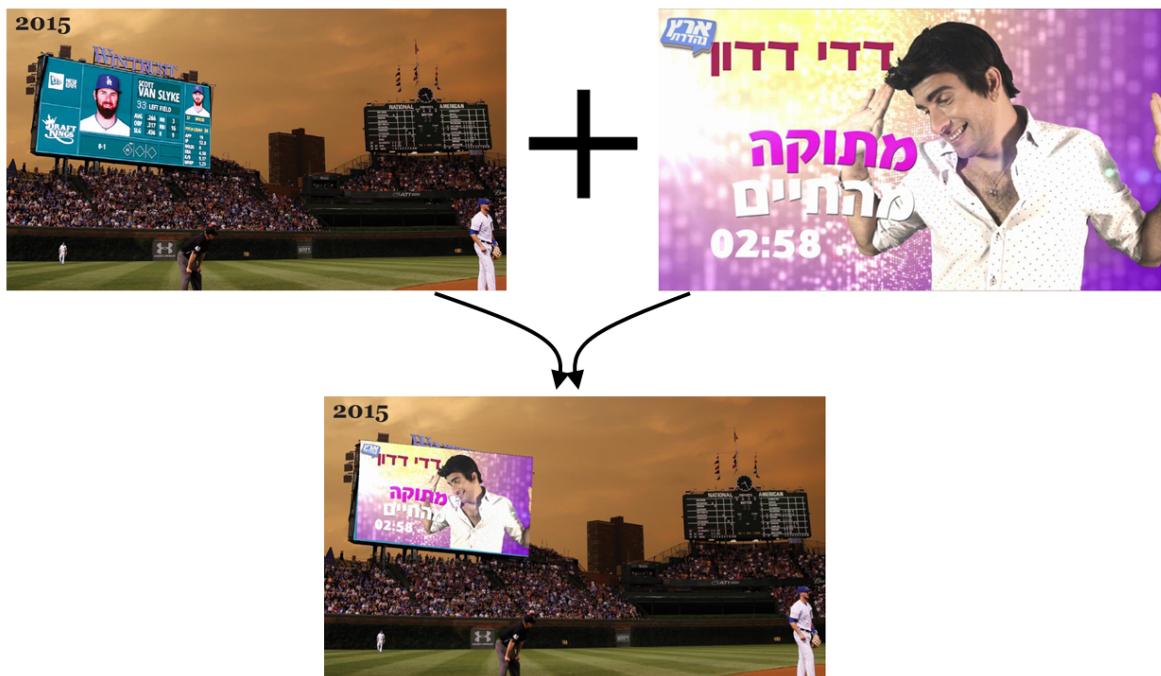
The RANSAC algorithm is extremely simple, **but it often**

- Does not produce correct model with user-defined probability
- Outputs an inaccurate model
- Does not handle degeneracies
- Can be sped up (by orders of magnitude)
- Does not guarantee minimum running time
- Needs information about scale of the noise
- Does not handle multiple models efficiently

Many improved algorithms:

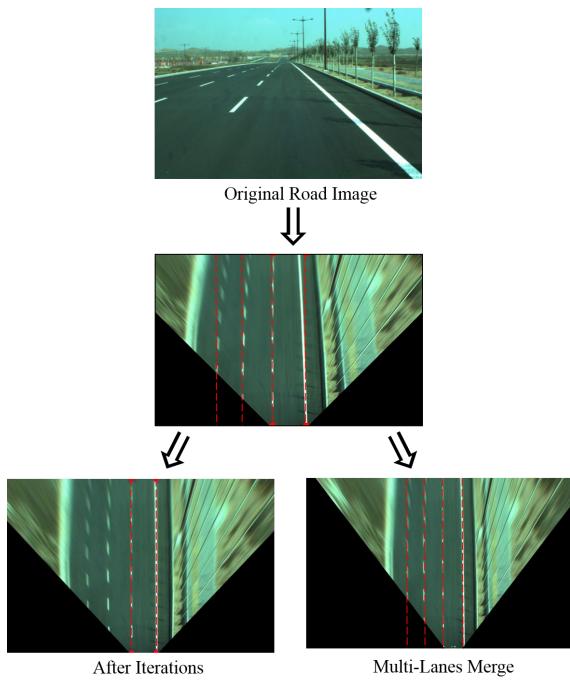
- [PROSAC](#)
 - Key idea is to use to assume that the similarity measure predicts correctness of a match
- [Randomized RANSAC](#)
 - Each step take a random subset of the query points and perform RANSAC
- [KALMANSAC](#)
- [and more...](#)
- Estimating homography with RANSAC in OpenCV: `cv2.findHomography(src_pts, dst_pts, cv2.RANSAC)`

Cool application 1: Planting images into other images



- result achieved with inverse warping and an affine transform

Cool application 2: BirdEye



- [Image Source](#)

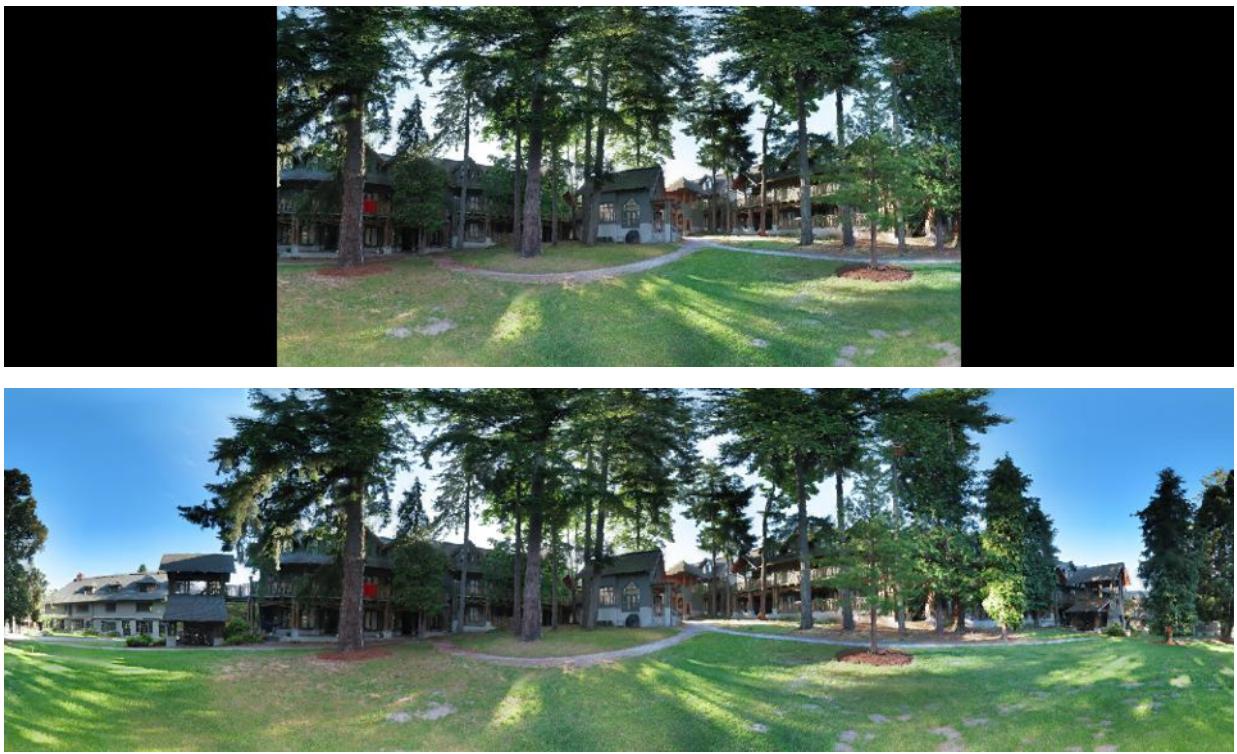
Panorama

- [Warping](#)
- [Image Blending \(Feathering\)](#)



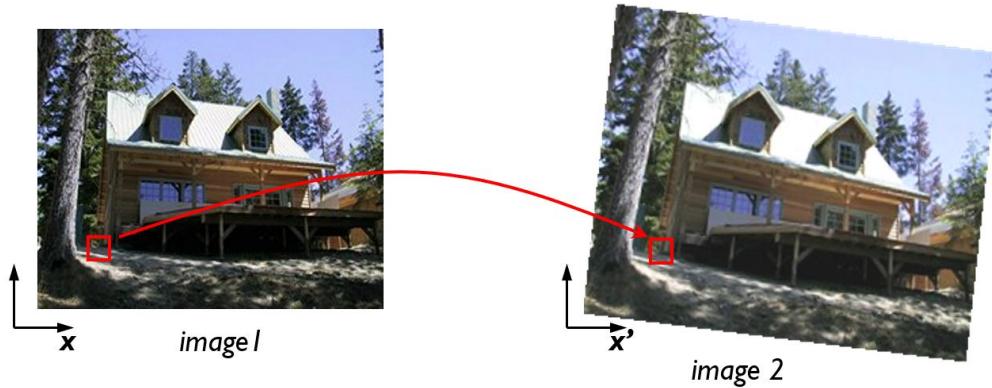
Obtain a wider angle view by combining multiple images





Warp - What we need to solve?

- Given source and target images, and the transformation between them, how do we align them?
- Send each pixel x in image1 to its corresponding location x' in image 2

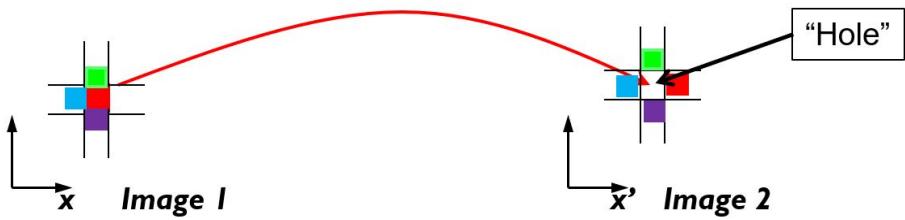


Forward Warping

- What if pixel lands “between” two pixels?
- Answer: add “contribution” to several pixels and normalize (splatting)

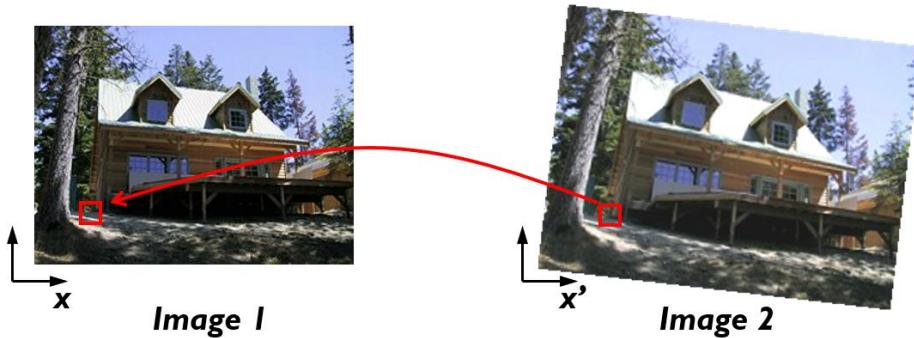


- Limitation: Holes (some pixels are never visited)

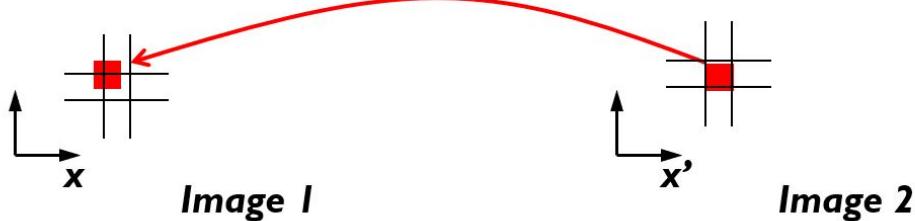


Inverse Warping

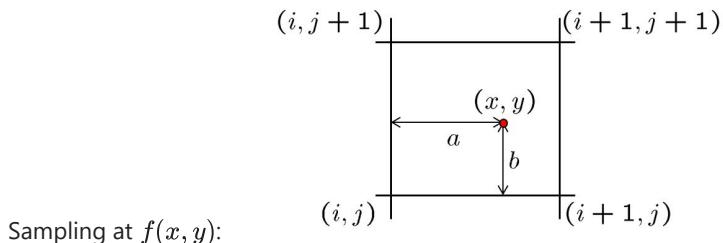
- For each pixel x' in image 2 find its origin x in image 1
- Problem: What if pixel comes from "between" two pixels?



- Answer: interpolate color value from neighbors



Bilinear Interpolation



$$\begin{aligned} f(x, y) = & (1 - a)(1 - b)f[i, j] \\ & + a(1 - b)f[i + 1, j] \\ & + abf[i + 1, j + 1] \\ & + (1 - a)bf[i, j + 1] \end{aligned}$$

Python:

- `interp2d()` - <https://docs.scipy.org/doc/scipy/reference/generated/scipy.interpolate.interp2d.html>

In [20]:

```
import numpy as np
from scipy import interpolate
import matplotlib.pyplot as plt

x = np.arange(-5.01, 5.01, 0.25)
y = np.arange(-5.01, 5.01, 0.25)
xx, yy = np.meshgrid(x, y)
z = np.sin(xx**2+yy**2)

## Cubic
f = interpolate.interp2d(x, y, z, kind='cubic')

xnew = np.arange(-5.01, 5.01, 1e-2)
ynew = np.arange(-5.01, 5.01, 1e-2)
znew = f(xnew, ynew)

plt.figure()
plt.subplot(121)
plt.imshow(z,vmin=-1,vmax=1)
plt.subplot(122)
plt.imshow(znew,vmin=-1,vmax=1)
plt.show()

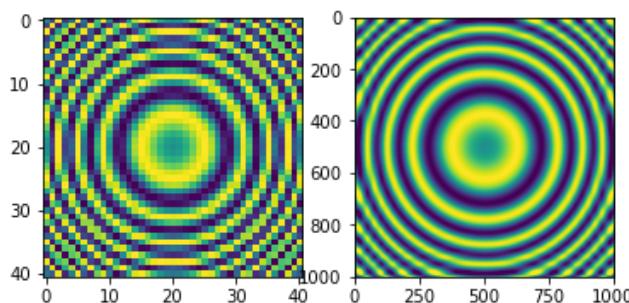
plt.figure()
plt.plot(x, z[0, :], 'ro-', xnew, znew[0, :], 'b-')
plt.show()

##Linear
f = interpolate.interp2d(x, y, z, kind='linear')

xnew = np.arange(-5.01, 5.01, 1e-2)
ynew = np.arange(-5.01, 5.01, 1e-2)
znew = f(xnew, ynew)

plt.figure()
plt.subplot(121)
plt.imshow(z,vmin=-1,vmax=1)
plt.subplot(122)
plt.imshow(znew,vmin=-1,vmax=1)
plt.show()

plt.figure()
plt.plot(x, z[0, :], 'ro-', xnew, znew[0, :], 'b-')
plt.show()
```



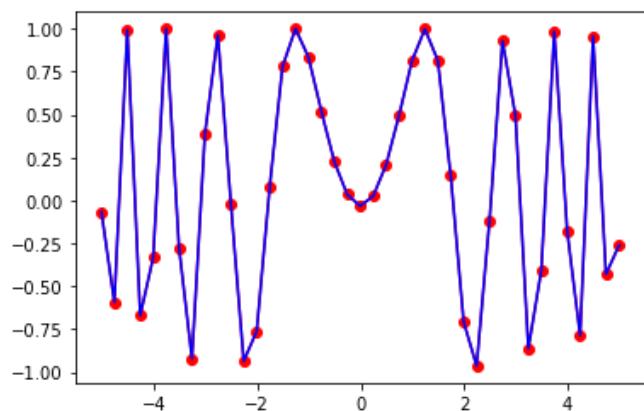
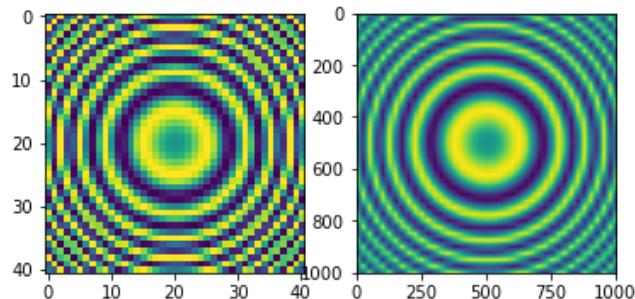
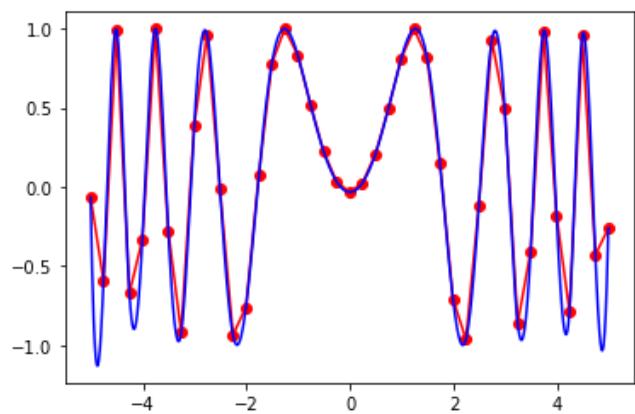
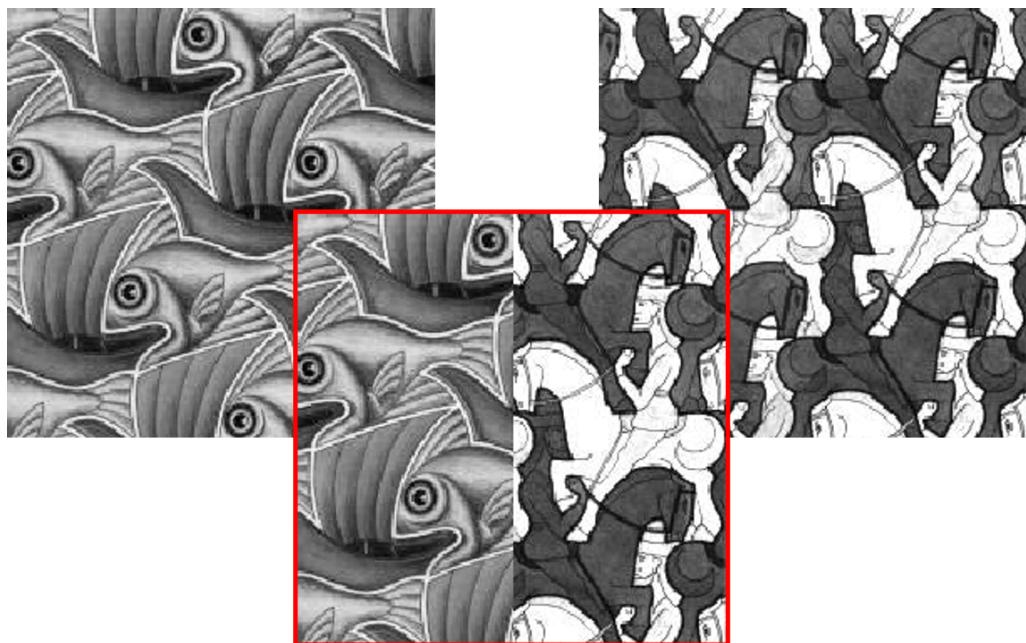


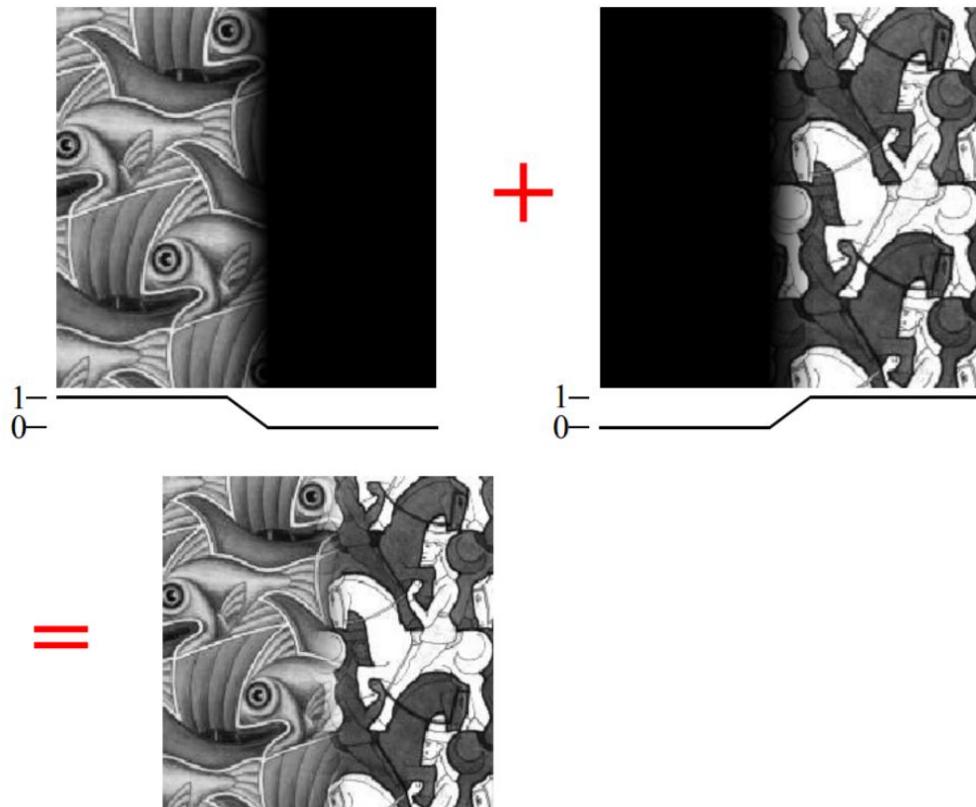
Image Blending



- Alpha blending

- Pyramid blending

Alpha Blending



Pyramid Blending:

1. Build a Gaussian pyramid for each image
2. Build the Laplacian pyramid for each image
3. Decide/find the blending border (in the example: left half belongs to image 1, and right half to image 2 -> the blending border is `cols/2`)
 - Split by index, or
 - Split using a 2 masks (can be weighted masks)
4. Construct a new mixed pyramid - mix each level separately according to (3)
5. Reconstruct a blend image from the mixed pyramid

In [21]:

```
def create_pyrs(A,B):
    # generate Gaussian pyramid for A
    G = A.copy()
    gpA = [G]
    for i in range(6):
        G = cv2.pyrDown(G)
        gpA.append(G)
    # generate Gaussian pyramid for B
    G = B.copy()
    gpB = [G]
    for i in range(6):
        G = cv2.pyrDown(G)
        gpB.append(G)
    # generate Laplacian Pyramid for A
    lpA = [gpA[5]]
    for i in range(5,0,-1):
        GE = cv2.pyrUp(gpA[i])
        L = cv2.subtract(gpA[i-1],GE)
        lpA.append(L)
    # generate Laplacian Pyramid for B
    lpB = [gpB[5]]
    for i in range(5,0,-1):
        GE = cv2.pyrUp(gpB[i])
        L = cv2.subtract(gpB[i-1],GE)
```

```

        lpB.append(L)
    return lpA,lpB

def blend_images(A,B):
    lpA,lpB = create_pyrs(A,B)
    # Now add left and right halves of images in each level
    LS = []
    for la,lb in zip(lpA,lpB):
        rows,cols,dpt = la.shape
        ls = np.hstack((la[:,0:int(cols/2)], lb[:,int(cols/2):])) #mixing can also be done with a mask
        LS.append(ls)
    # now reconstruct
    ls_ = LS[0]
    for i in range(1,6):
        ls_ = cv2.pyrUp(ls_)
        ls_ = cv2.add(ls_, LS[i])
    # image with direct connecting each half
    real = np.hstack((A[:,0:int(cols/2)],B[:,int(cols/2):]))
    return real, ls_

def switch_texture(A,B):
    lpA,lpB = create_pyrs(A,B)
    # Now add left and right halves of images in each level
    LS = []
    # for la,lb in zip(lpA,lpB):
    #     rows,cols,dpt = la.shape
    #     ls = np.hstack((la[:,0:int(cols/2)], lb[:,int(cols/2):])) #mixing can also be done with a ma
    #     LS.append(ls)
    # now reconstruct
    ls_ = lpA[0]
    for i in range(1,6):
        ls_ = cv2.pyrUp(ls_)
        ls_ = cv2.add(ls_, lpB[i])
    # image with direct connecting each half
    real = np.hstack((A[:,0:int(cols/2)],B[:,int(cols/2):]))
    return real, ls_

def alpha_blend(A,B,MASK):
    A = np.float32(A)
    B = np.float32(B)
    return np.uint8(A*MASK), np.uint8(A*MASK+B*(1-MASK))

```

In [22]:

```

A = cv2.imread('../assets/apple.jpg')
B = cv2.imread('../assets/orange.jpg')
real,ls_ = blend_images(A,B)

ls_ = cv2.cvtColor(ls_,cv2.COLOR_BGR2RGB)
real = cv2.cvtColor(real,cv2.COLOR_BGR2RGB)

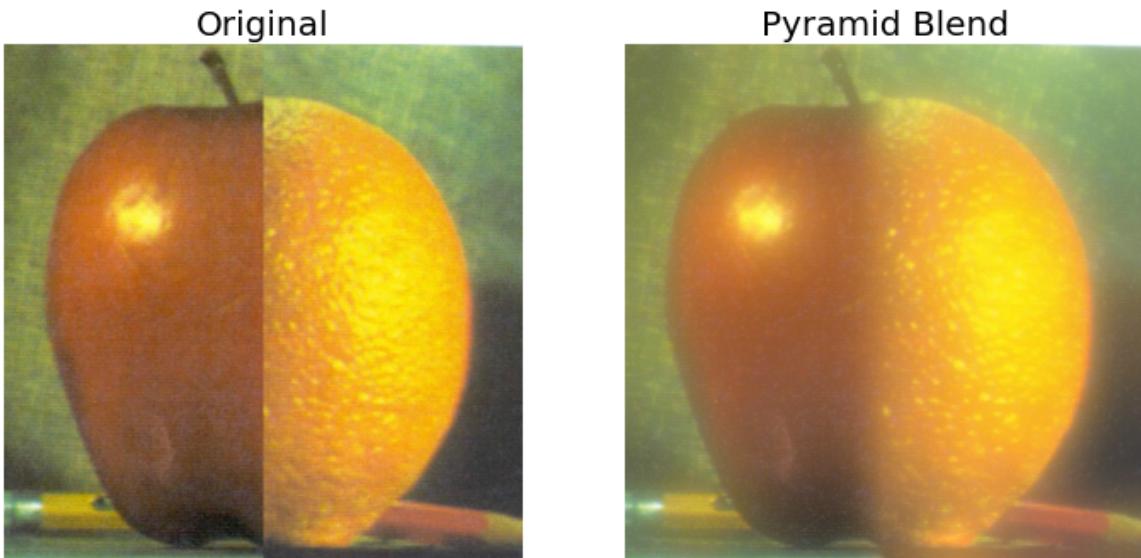
```

In [23]:

```

plot_images([real,ls_],['Original','Pyramid Blend'],(1,2),figsize=(12,12),fontsize=20)

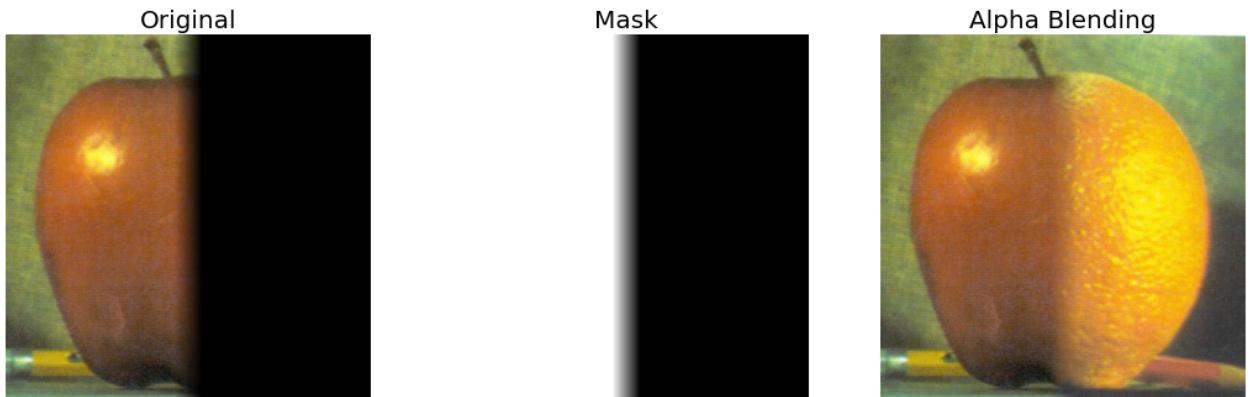
```



- Alpha blending example:

```
In [24]: MASK = np.ones_like(A,np.float32)
rows,cols,dpt = MASK.shape
w=20
v_dec = np.linspace(1,0,2*w)
MASK[:,int(cols/2):]=0
MASK[:,(int(cols/2)-w):(int(cols/2)+w)]=np.tile(np.reshape(v_dec,[1,-1,1]),[rows,1,3])
real,ls_ = alpha_blend(A,B,MASK)
ls_ = cv2.cvtColor(ls_,cv2.COLOR_BGR2RGB)
real = cv2.cvtColor(real,cv2.COLOR_BGR2RGB)
```

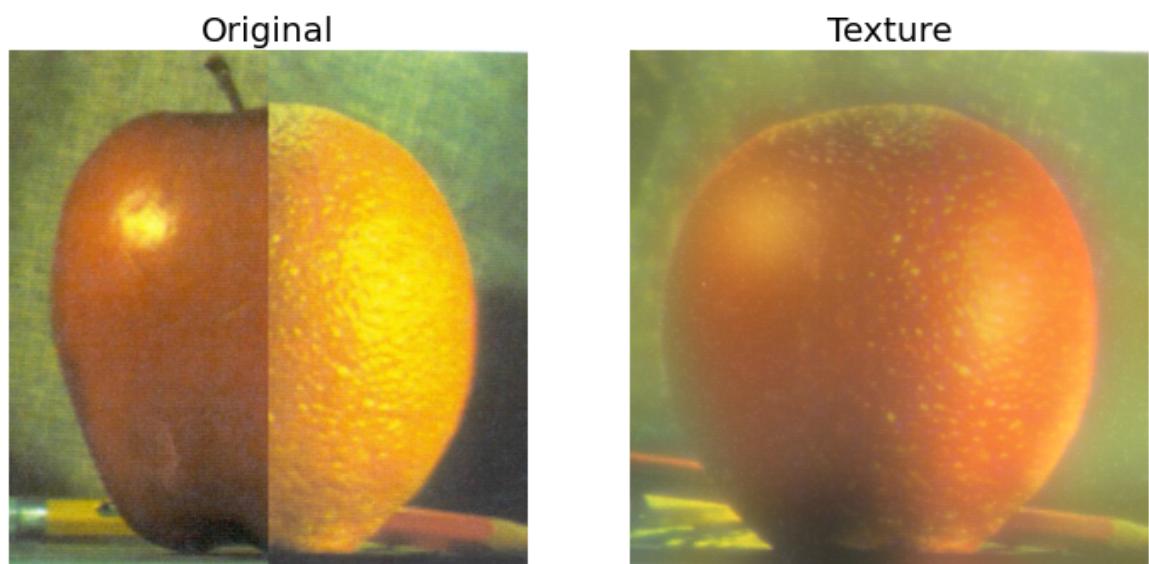
```
In [25]: plot_images([real,MASK,ls_],['Original','Mask','Alpha Blending'],(1,3),figsize=(18,12),fontsize=20)
```



- Stylize image using pyramids:

```
In [26]: real,ls_ = switch_texture(A,B)
ls_ = cv2.cvtColor(ls_,cv2.COLOR_BGR2RGB)
real = cv2.cvtColor(real,cv2.COLOR_BGR2RGB)
```

```
In [27]: plot_images([real,ls_],['Original','Texture'],(1,2),figsize=(12,12),fontsize=20)
```



Blending Improvements

- Many algorithms have different variations of combining alpha and pyramid blending (different masks for different frequencies)
- Find the boundaries using **segmentation**



Panorama - Summary

- Detect features
- Compute transformations between pairs of frames
- Can Refine transformations using RANSAC
- Warp all images onto a single coordinate system
- Find mixing borders (e.g. using segmentation)
- Blend



Transformations in Deep Learning

- Can we incorporate transformations in the pipeline of a deep learning algorithm?
 - Moreover, can we accelerate these transformations by performing them on a GPU?
- **YES!**

Kornia - Computer Vision Library for PyTorch



- Kornia is a differentiable computer vision library for PyTorch
 - That means you can have gradients for the transformations!
- Inspired by OpenCV, this library is composed by a subset of packages containing operators that can be inserted within neural networks to train models to perform image transformations, epipolar geometry, depth estimation, and low-level image processing such as filtering and edge detection that operate directly on tensors.
- Check out `kornia.geometry` - <https://kornia.readthedocs.io/en/latest/geometry.html>
- Warp image using perspective transform

```
import torch
import kornia

frame: torch.Tensor = load_video_frame(...)

out: torch.Tensor = (
    kornia.rgb_to_grayscale(frame)
)
```



Recommended Videos



Warning!

- These videos do not replace the lectures and tutorials.

- Please use these to get a better understanding of the material, and not as an alternative to the written material.

Video By Subject

- Homography
 - Image geometry and planar homography - ENB339 lecture 9: Image geometry and planar homography
 - Homography - [Homography in computer vision explained](#)
- Transformations - [Lect. 5\(1\) - Linear and affine transformations](#)
- Matching Local Features
 - SIFT - [CSCI 512 - Lecture 12-1 SIFT](#)
- RANSAC (see [Tutorial 2](#))



Credits

- EE 046746 Spring 2020 - [Dahlia Urbach](#)
- Slides - Elad Osherov (Technion), Simon Lucey (CMU)
- Multiple View Geometry in Computer Vision - Hartley and Zisserman - Section 2
- [Least-squares Solution of Homogeneous Equations](#) - Center for Machine Perception - Tomas Svoboda
- [Computer vision: models, learning and inference](#), Simon J.D. Prince - Section 15.1
- [Computer Vision: Algorithms and Applications](#) - Richard Szeliski - Sections 2,4,6, 9 ([Free for Technion students via remote library](#))
- Icons from [Icon8.com](#) - <https://icon8.com>